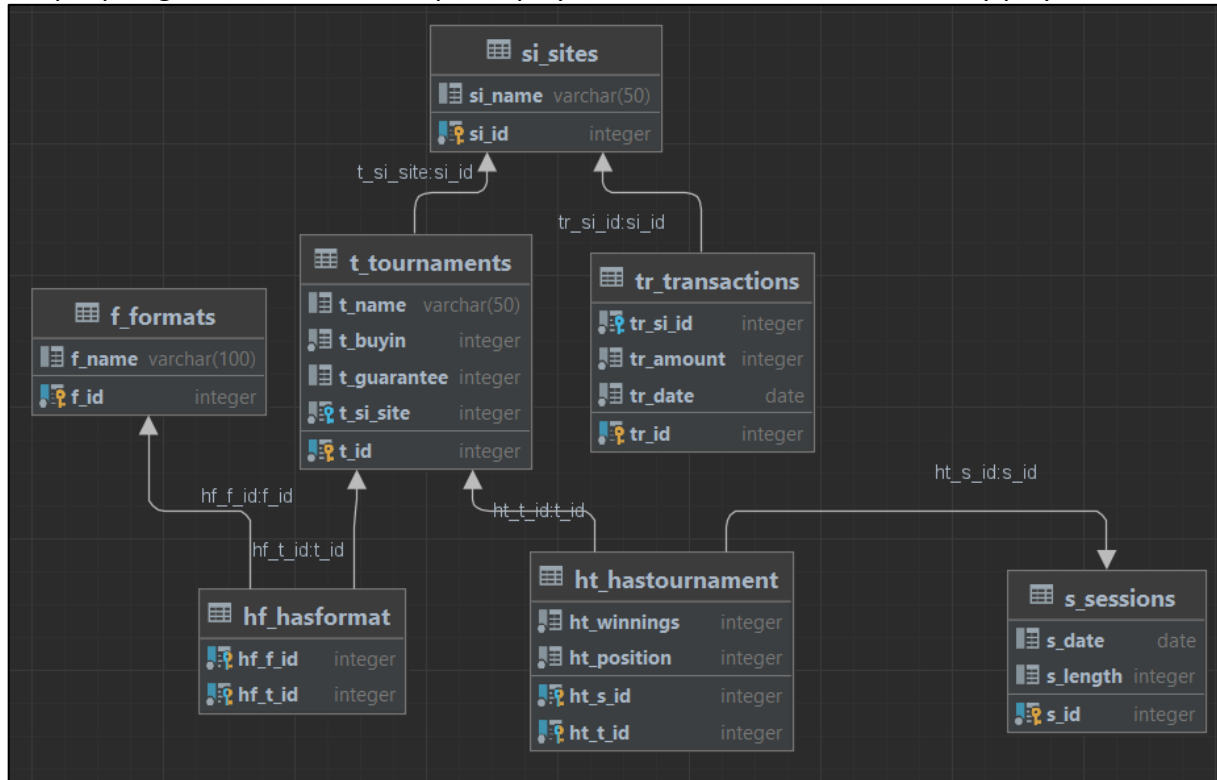


## Schritt 1

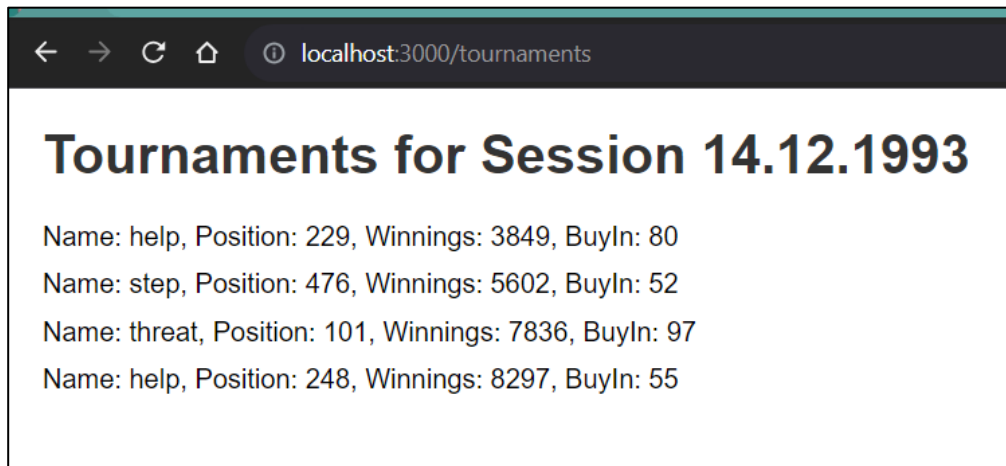
Simple postgres-db for an online poker player to track the tournaments they play:



For the frontend we have the functionality that the user can select the date of a session from a drop down to then see the tournaments they played on that day and the results they achieved:

The screenshot shows a web browser window at localhost:3000 with the title "Sessions by Date". Below the title, there is a label "Select ID:" followed by a dropdown menu. The dropdown menu is open, displaying a list of dates from 6.12.1993 to 31.12.1993. The date 14.12.1993 is currently selected and highlighted in blue. To the right of the dropdown is a blue button labeled "Get Sessions".

Pressing Get Sessions leads to:



## Schritt 2

Model can be found [here](#). To fulfill all the criteria, we wrote a runnable implementation like Mr. Schletz's generating data [here](#).

Sample data in DB:


```
_id: ObjectId('657c64857b956c7ba6f1a1e7')
date: 2005-05-15T00:00:00.000+00:00
length: 2
▼ stats: Array (5)
  ▼ 0: Object
    winnings: 8072
    position: 1664
    ▼ tournament: Object
      name: "treat"
      buyin: 13
      guarantee: 9483
    ▼ site: Object
      name: "back"
    ▼ formats: Array (2)
      ▼ 0: Object
        name: "early"
      ▼ 1: Object
        name: "meeting"
  ▶ 1: Object
  ▶ 2: Object
  ▶ 3: Object
  ▶ 4: Object
```

## Schritt 3


Runnable with the Jupyter Notebook PerformanceComparison.ipynb

### Write Test

We insert 10, 100, 1000, 10000 sessions, each with 1-5 tournaments. These inserts are the basis for all future tests.

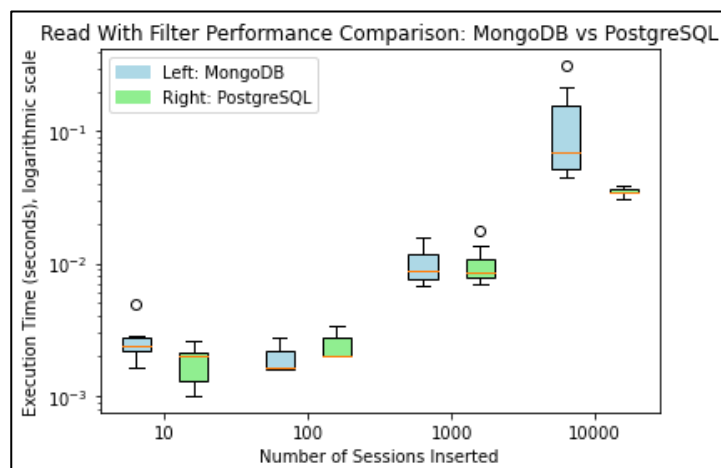


[Read All](#)




[Read With Filter](#)

We Filter by the session length, selecting all sessions with the length of 1, which should be 20%



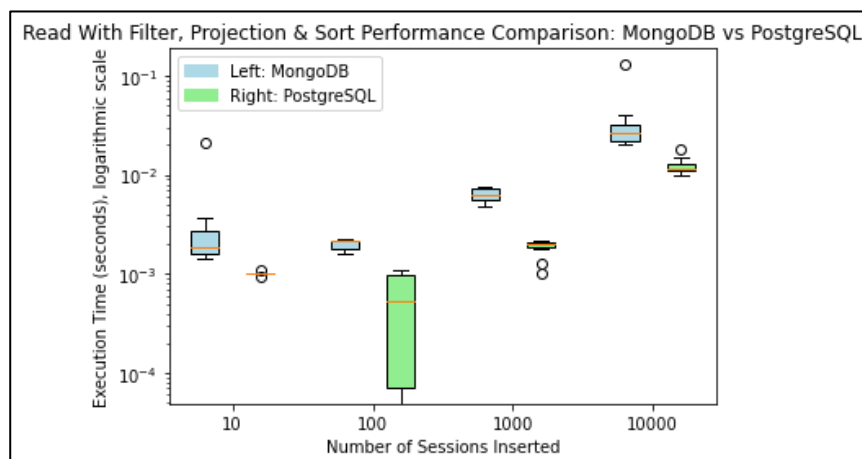
### Read With Filter & Projection

We project so that we only select a subset of the data provided in sessions and tournaments. Due to this we also use fewer joins in the sql implementation as we are not getting all the data.




### Read With Filter, Projection & Sorting

We add sorting by the date.



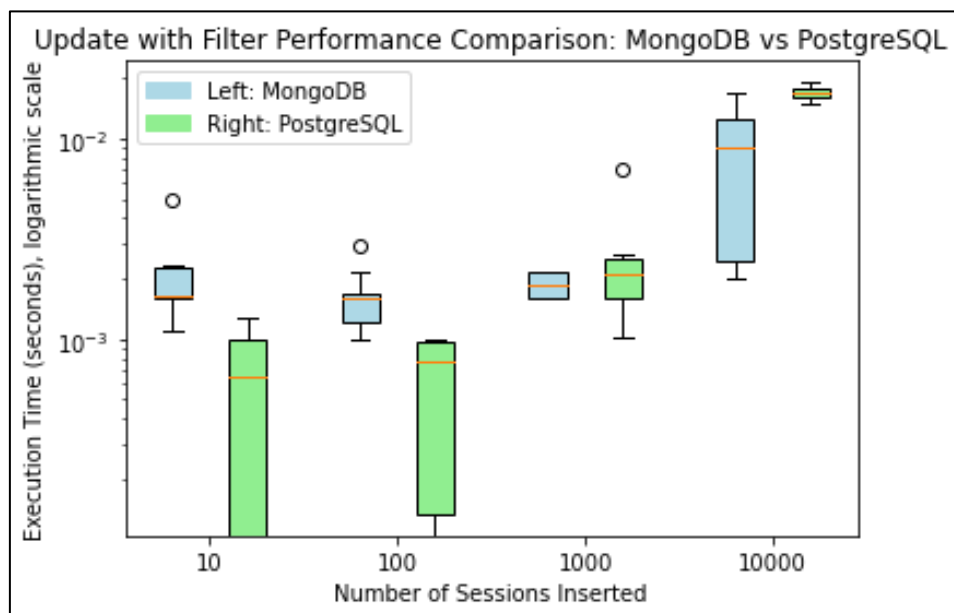
### Update All

We increment every length of the session by one.




### Update With Filter

We only update the sessions with a length of 1, approximately 20%.



### Delete All



## Bonus

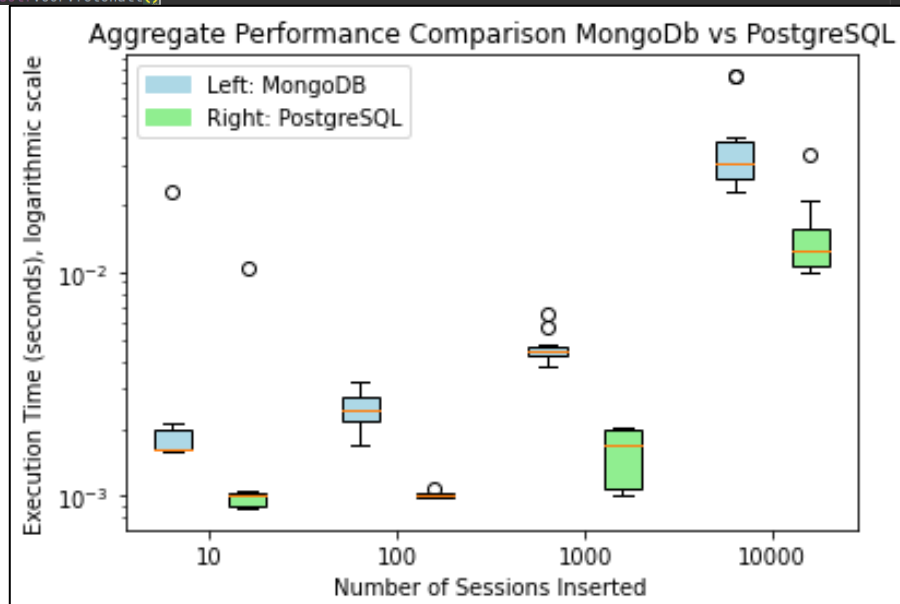
### 1. Aggregation

Mongo aggregate:

```
def aggregate(self):
    return list(self.mycol.aggregate([{"$unwind": "$stats"}, {"$group": {"_id": "$stats.tournament.name", "total": {"$sum": "$stats.winnings"}}}]))
```

SQL version:

```
def group_by(self):
    self.cur.execute(f"SELECT t_name, SUM(ht_winnings) FROM t_tournaments t JOIN ht_hasTournament ht ON t.t_id = ht.ht_t_id GROUP BY t_name;")
    return self.cur.fetchall()
```



### 2. Referencing

Model

The adapted model can be found [here](#). Compared to the [old one](#) there isn't too much change, except that N:M can now be modelled properly. The data faker for referencing can be found [here](#). This uses the bson DBref feature:

```
tournament_documents.append({
    'name': tournament_data[0],
    'buyin': tournament_data[1],
    'guarantee': tournament_data[2],
    'site': DBRef(collection='sitesRef', id=site_instance),
    'formats': single_format_instances
})
```

In the new model the session looks like this:

```
_id: ObjectId('657c6b2b7b956c7ba6f28d80')
date: 2004-03-07T00:00:00.000+00:00
length: 3
```


While then the session stats collection works with references:

```
_id: ObjectId('657c6ce37b956c7ba6f351ac')
winnings: 726
position: 211
session: DBRef('sessionsRef', '657c6cd67b956c7ba6f32a9c')
tournament: DBRef('tournamentsRef', '657c6cd67b956c7ba6f32a65')
```

**NOTE:** We are using **pyMongo** which is not a proper ORM like **mongoengine**. This means that there is no domain model-level support for DBRef. Luckily there is still support for manual db-referencing when making inserts, which we will use.

### Performance Comparison

The use of references leads to performance slowdowns across the board. One thing to note is that the inserts are not 100% comparable, as with the base mongo implementation the data is first prepared in document form and only the insertion time is measured. For the referencing, however, this was not feasible and there is likely a bit of overhead in those numbers which is not directly attributed to the database performance.

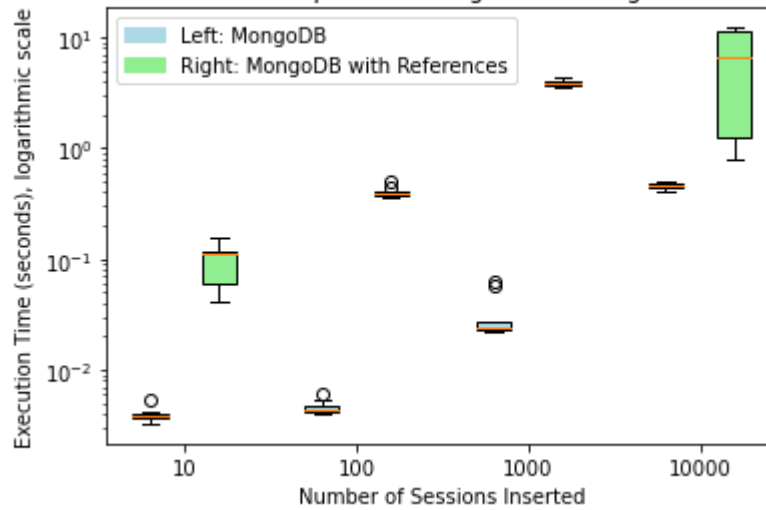


For the queries we used the lookup operation:


```
def find_all(self):
    pipeline = [
        {"$lookup": {
            "from": "tournamentsRef", # Replace with the actual name of the tournaments collection
            "localField": "stats.tournament.$id",
            "foreignField": "_id",
            "as": "tournament"
        }},
        {"$lookup": {
            "from": "sessionsRef", # Replace with the actual name of the sessions collection
            "localField": "stats.session.$id",
            "foreignField": "_id",
            "as": "session"
        }},
        {"$project": {"_id": 0, "date": 1, "length": 1, "stats.winnings": 1, "tournament.name": 1, "session.length": 1}}
    ]

    result = list(self.mycol_sessions.aggregate(pipeline))
    return result
```

Read All Performance Comparison: MongoDB vs MongoDB with References



Read With Filter Performance Comparison: MongoDB vs MongoDB with References

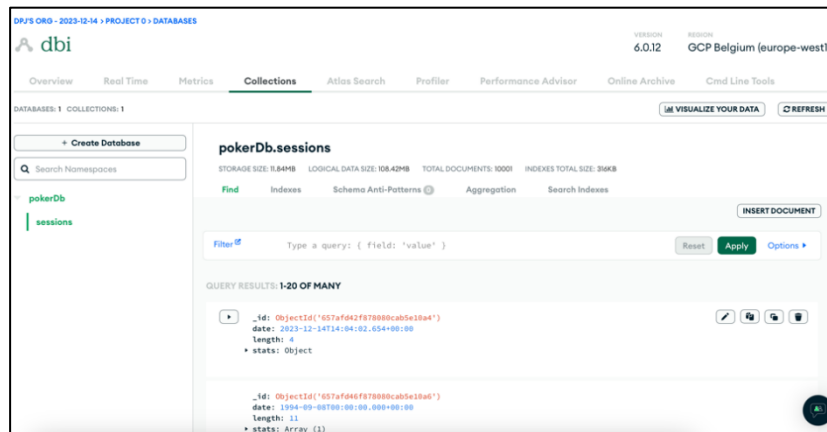


### 3. Cloud

#### General

We did give Atlas Cloud a try and it works amazingly:

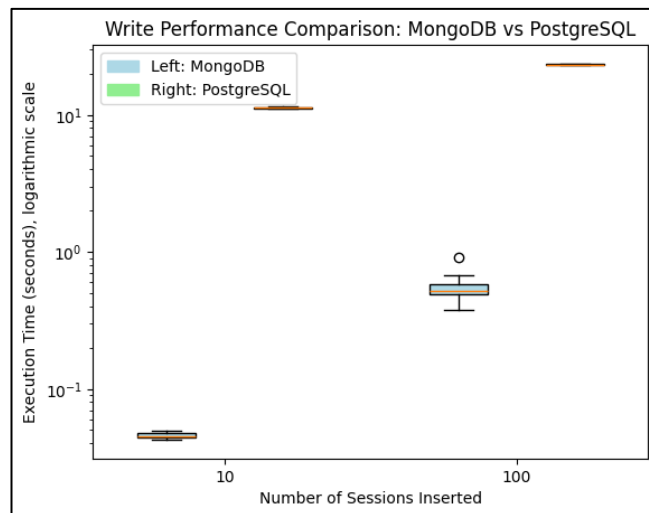


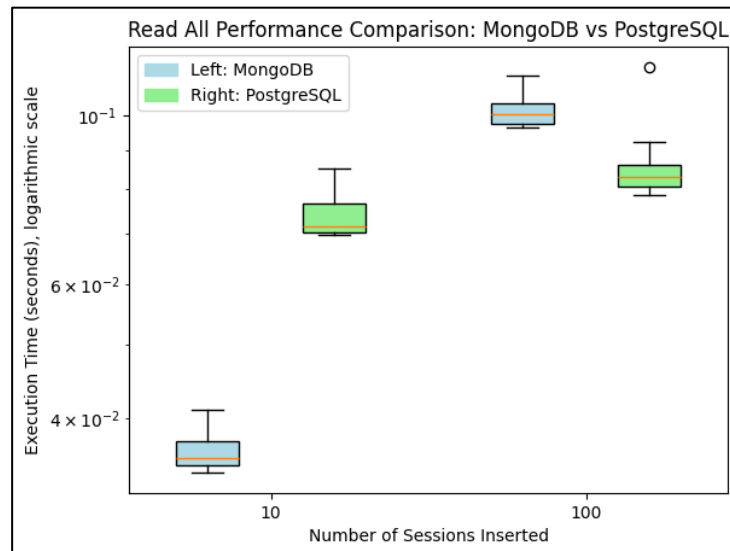


For the relational DB we opted not to use Azure because we didn't have credits left. We therefore rely on Google Cloud. More details on the infrastructure are to be found in the [Something Cool](#) section of this document.

### Benchmark

(note that our google cloud instance does have very limited compute resources; especially when compared to a database as a service solution like Atlas. Also, the network forwarding is a lot more complex – all this impacts the results – so take with a big grain of salt):





#### 4. Mongo Frontend

Mongo Frontend can be viewed in [here](#). As requested, we implemented the **filter feature**. As discussed with you, because this is the frontend-optimized DB version the filtering is done on the frontend as it is provided with all the sessions. The drawbacks of this approach are obvious, more data more loading time, but for frontend dev experience this is the simplest, easiest/most optimized.

Sessions by Date

Select Session: 11/7/2001

_id	6579cbd8f878080cab5de9a2														
date	2001-11-07T00:00:00.000Z														
length	11														
winnings	4725														
position	133														
formats	<table border="1"><tr><td>name</td><td>final</td></tr><tr><td>buyin</td><td>71</td></tr><tr><td>guarantee</td><td>5732</td></tr><tr><td>site</td><td>name 3</td></tr><tr><td>0</td><td><table border="1"><tr><td>name</td><td>2</td></tr><tr><td>1</td><td>name 5</td></tr></table></td></tr></table>	name	final	buyin	71	guarantee	5732	site	name 3	0	<table border="1"><tr><td>name</td><td>2</td></tr><tr><td>1</td><td>name 5</td></tr></table>	name	2	1	name 5
name	final														
buyin	71														
guarantee	5732														
site	name 3														
0	<table border="1"><tr><td>name</td><td>2</td></tr><tr><td>1</td><td>name 5</td></tr></table>	name	2	1	name 5										
name	2														
1	name 5														

#### 5. Index Runtime

Since in our testing we filtered by length we decided to create an index on the length. This was done using the MongoDB compass app:

## Create Index

pokerDb.sessions

Index fields

length 1 (asc) +

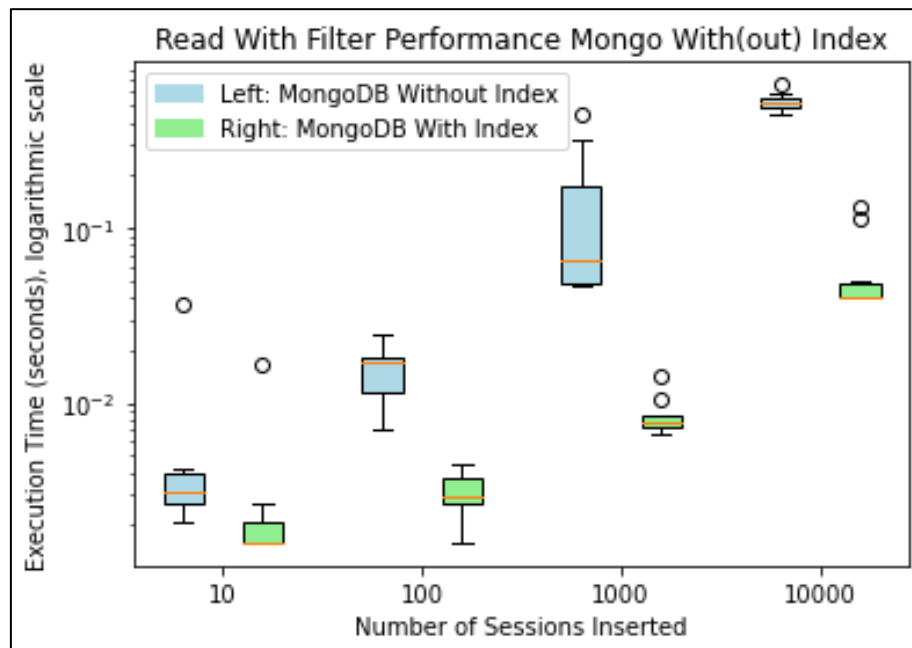
> Options

Cancel
Create Index

Name and Definition	Type	Size	Usage	Properties
> _id_	REGULAR ⓘ	249.9 KB	14 (since Thu Dec 14 2023)	UNIQUE ⓘ
> length_1	REGULAR ⓘ	69.6 KB	0 (since Thu Dec 14 2023)	

### Test

Testing was again done in the Jupyter notebook. The results of a run without indexing were saved and compared to those after indexing. An increase in performance can be seen.



## 6. Something Cool

We use Google Cloud to run our relational database. To make things **even cooler** we don't simply run the DB as a service but rather in Kubernetes environment as a full deployment! The deployment files can be found [here](#).

Our cluster:

NAME	LOCATION	MASTER_VERSION	MASTER_IP	MACHINE_TYPE	NODE_VERSION	NUM_NODES	STATUS
schmangie	eu-west1-b	1.28.3-gke.1203001	35.195.228.76	n1-standard-2	1.28.3-gke.1203001	2	RUNNING

